



Code Review MAS - DAILY PLAN

◆ Day 1 – MAS Design + Mesa Setup

✓ Goals:

- Define system architecture (Coder, Reviewer, later Architect)
- Install Mesa, structure model + scheduler

Tasks:

- Create project folder (`code_review_mas`)
 - Install Mesa + basic dependencies
 - Define `CoderAgent`, `ReviewerAgent` classes
 - Create `CodeReviewModel` with `run_task()` logic
-

◆ Day 2 – Instrument Agent Behavior

✓ Goals:

- Add OpenTelemetry spans to agent `step()` methods
- Export traces to stdout or local file

Tasks:

- Install OpenTelemetry SDK (`opentelemetry-sdk`)
 - Use `tracer.start_as_current_span(...)` for:
 - Task implementation
 - Code review
 - Add span attributes (task ID, output, errors)
-

◆ Day 3 – Add Metrics + Internal Testing

✓ Goals:

- Add internal metrics like code/task similarity
- Inject synthetic errors

Tasks:

- Use `SentenceTransformer` or `sklearn` to compute cosine similarity between task + generated code
 - Simulate edge cases (ambiguous task, bad code)
 - Output final span with `{"similarity": 0.81, "errors": 1}`
-

◆ Day 4 – LLM Task Generation

✓ Goals:

- Integrate GPT-4-turbo for `generate_task()` (OpenAI API)

Tasks:

- Set up API key, use `openai.ChatCompletion.create(...)`
 - Prompt: “Give me a backend feature request”
 - Extract output, send to `CodeReviewModel.run_task()`
-

◆ Day 5 – LLM-based Workflow Planning

✓ Goals:

- Decompose LLM task into subtasks (future “PlannerAgent” role)

Tasks:

- Prompt LLM to create multi-step plan for ambiguous input
 - Assign subtasks to coder/reviewer manually
 - Store full trace of LLM plan + MAS execution
-

◆ Day 6 – Stress Test the Pipeline

✓ Goals:

- Run 20–50 tasks through system, verify trace span richness

Tasks:

- Loop: `for task in generate_tasks(n=50): model.run_task(task)`
 - Save results as JSONL or Pandas DataFrame
 - Check span coverage: input → plan → code → review → outcome
-

◆ Day 7 – OpenTelemetry Export to Jaeger

✓ Goals:

- Setup Jaeger locally, send spans from MAS to UI

Tasks:

- Install Jaeger via Docker (`docker run ...`)
 - Add `OTLPSpanExporter` to your tracing pipeline
 - Visualize agent timelines in Jaeger (span tree per task)
-

◆ Day 8 – Add Architect Agent (Optional)

✓ Goals:

- Architect injects design constraints (e.g., “optimize for speed” vs “secure login”)

Tasks:

- Create `ArchitectAgent`
 - Feed contradictory goals into task
 - Measure downstream error rate or agent conflict
-

◆ Day 9–10 – MAST-style Analysis

✓ Goals:

- Apply metrics to trace logs:
 - Misalignment scores
 - Error propagation
 - Role-based span timelines

Tasks:

- Parse OTLP data (or span logs) into Pandas
 - Cluster interactions with high error/similarity divergence
 - Visualize agent interactions (networkx, Sankey, Seaborn)
-

◆ Day 11–12 – Compare Experiments

✓ Goals:

- Compare:
 - With vs. without Architect

- Ambiguous vs. precise task prompts
- Reviewer strictness levels

Tasks:

- Run comparative batches
- Aggregate metrics (mean error, avg span length, task/code alignment)
- Save visuals and notebook summaries

◆ Day 13 – Report & Visualizations

✓ Goals:

- Write markdown/PDF report (GitHub/Overleaf)
- Include code, spans, architecture diagram

Tasks:

- Document architecture, span schema, LLM prompts
- Add Jaeger screenshots or span timelines
- Summarize insights (e.g., “error rate rises with ambiguity”)

◆ Day 14 – Final Polish

✓ Goals:

- Clean up repo
- Optional: submit as research artifact or internal demo

Tasks:

- Push code + data to GitHub
- Tag LLM prompts + responses
- Add README with system overview and run instructions